Resource-constraint project scheduling with task preemption and setup times by heuristically augmented SAT solver

Jasper Vermeulen[1]
Supervisor: Emir Demirović[1]

[1]EEMCS, Delft University of Technology, The Netherlands
J.Vermeulen-1@student.tudelft.nl

Abstract

This will only be written once all experiments have done and the conclusion is on paper.

1 Introduction

The problem of scheduling tasks arises in industries all the time. It is not hard to imagine that generating an optimised schedule can be of great profit for production or logistic operations. For example, optimisation can minimise the overal required time or minimise the delay before starting a task. Because this type of problem is so prevalent it has already been subject to much research.

Formally this specific type of problem is known as the resource-constrained project scheduling problem (RCPSP). The standard RCPSP has a set of tasks that each require a specified resource for a given amount of time and precedence restrictions are given for the tasks. On its own, this problem definition for standard RCPSP is limited and of little use to realistic applications where additional constraints must be followed or more options are available. To make sure the researched algorithms solving the scheduling problem would have a wider use case, many variations and extensions to the problem definition have been classified over time [1], [2]. More recently, the variations and extensions have also been surveyed and put into a structured overview [3].

For this research, the preemptive resource-constrained project scheduling problem with setup times (PRCPSP-ST) variant is under study. Preemption allows a task to be interrupted during its scheduled time by another task. Each interruption can be seen as splitting the task into multiple smaller activities. The setup times are introduced for each interruption in an task to discourage endless splits resulting in a chaotic schedule. Both a model for allowing preemption [4] and including setup times [5] have already been established. Both models have been combined and a proposed algorithm for it was found to result in a reduction of the makespan compared to the optimal schedule without task preemption [6]. Within this algorithm, the activities are split into all possible integer time segments and a SAT solver makes a selection from these segments [7]. The resulting list was used to construct a schedule with a genetic algorithm established in earlier research [8].

In the research done on solving RCPSP variants, the focus has been on heuristic and meta-heuristic algorithms. These algorithms are usually variants of branch-and-bound algorithm [4] or a form of genetic algorithms [9] that were established for the standard RCPSP.

SAT solvers are a very general tool they can be used on any algorithmic problem as long as it is encoded as the required input for the solver. SAT solvers have been used as a part of the algorithm but using a SAT solver as a complete solution to try and solve PRCPSP-ST instances has not been researched before. Heuristic algorithms by comparison are specialised by knowledge or an insight into the problem variation translated into a heuristic rule for the algorithm to use. Because the RCPSP is known to be strongly NP-hard [10], the SAT solver might not be more efficient than the heuristic and meta-heuristic at first. But, a SAT solver can also be modified to include a heuristic for the PRCPSP-ST variant. In a way, the addition of a heuristic can be seen as an augmentation of the SAT solver. The augmentation is making the solver more specialised for the PRCPSP-ST variant and could result in a better performance than other algorithms.

Because there is room to try and find out if a SAT solver can outperform the heuristic and meta-heuristic algorithms the main research question is: *Can the addition of a simple heuristic to a SAT solver algorithm used to solve to PRCP-SPST models reduce the average makespan of the resulting schedule when run for an equal amount time?*. The expectation for this research is to first show that a SAT solver can be used to solve PRCPSP-ST instances. And next, show that the heuristically augmented version of the SAT solver algorithm will result in a lower makespan than a heuristic algorithm when running an equal amount of time.

Report section structure goes here

2 Problem Formulation

The resource-constrained project scheduling problem is a strongly NP-hard algorithmic problem [10] with the objective is to minimize makespan (overall required time to finish all tasks).

2.1 RCPSP

RCPSP is about a project consisting of a set of tasks N which all have to completed to finish the project. Each task i has a duration d_i and a requires an amount $r_{i,k}$ of a resource type

k. A project provides a set of limited resources types R to process tasks each with an availability a_k constant throughout the project horizon. Tasks can be scheduled in timeslots as long as the overall resource type requirement does not exceed the provided amount at any time. Furthermore, a (possible empty) set of task pairs defines precedence relations A. The task pairs have a finish-start type precedence meaning that a task must be completed entirely before its successor can be started. Some additional assumptions are that each resource type required by any task is provided $k \in R$ and no single task will require more of a resource type than provided $r_{i,k} \leq a_k$.

This project structure can be modelled as an activity-on-the-node network (where activity also means task) $G = (N, A)$. The network is extended with 2 dummy tasks that model the start and finish of the project. These dummy tasks have a duration of 0 and no resource requirement. The makespan can now be difined as the starting time of the finish dummy task.

2.2 PRCPSP-ST

The RCPSP definition can be extended in different ways including task preemption and setup times [3], [11].

Preemption allows and activity to be paused after it has been started by the project. A preempted task in a schedule is like multiple individual tasks that each represent a segment of the original task. Preemption is only allowed at integer points of the task duration. In reality tasks might be preempted at any fraction of the duration but the infinite ways to split a task makes an algorithm much harder to define. A solution to approximate fractional preemption is rescaling the time units used project. When hours are scaled down to minutes for example, a task can be preempted on each minute instead of on the hour approaching a possible required granularity.

Setup time s is introduced as a way to try and prevent task being split into many impractical segments. This prevention is done by adding additional processing time (setup time) to task segments that start after a previous segment has been preempted. During the setup time the same amount and type of resource are required as the task itself. By penalizing preemption in this way algorithms will only introduce split tasks when the makespan can be improved in a meaningful way.

3 Heuristics and augmentation of SAT solvers

To evaluate the performance of a heuristically augmented SAT solver on PRCPSP-ST instances comparative data needs to be generated. To make sure the variables between algorithms are limited they were all designed, implemented and run on the same hardware for this research. This way the at least coding skills and hardware improvements over time are not introduced in the data. Improved performance of new algorithms only as a result of the use of more recent hardware has shown up in the past as described in section 5. Three algorithms are used to solve the same problem instances for an equal amount of time: a heuristic algorithm, SAT solver, heuristically augmented SAT solver.

The heuristic algorithm is an adapted version of the iterated greedy algorithm [12]. It was designed for flow-shop scheduling but with a few tweaks it can also be applied to RCPSP.

For a SAT solver to be used the PRCPSP-ST has to be encoded into conjunctive normal form boolean logic. When the encoding is made any SAT solver is able to provide feasible schedules. Because there is a clear objective to reduce the makespan a more advanced MAX-SAT solver is used. This will create feasible schedules while also trying to optimize to an objective function.

To augment a SAT solver for a specific problem the code of the solver has to be changed. In the case of this research a heuristic function will be added that improves the selection of a possible variable. The heuristic function will use knowledge about PRCPSP-ST to try and select a variable that will reduce the resulting schedule makespan.

3.1 Heuristic

As a heuristic solution a tweaked version of the iterated greedy algorithm is implemented [12]. This algorithm requires an activity list representation of the project. It start with a setup of an initial schedule and then iterates over a destruction phase and a construction phase until a time limit or number of iterations limit is reached.

An activity list representation allows a serial generation scheme to construct a feasible schedule [13]. The activity list represent a project as permutation vector of all the tasks. It is required that no task appears in the list after any of its successors. The serial generation scheme schedules all the task in the order of the list at the earliest possible start time that does not break precedence or resource restrictions. Because the generation scheme uses the tasks in order of the list tasks close to the front of the list can be seen as having a higher priority and are scheduled sooner by the algorithm. When the algorithm has scheduled all tasks the result is a left-justified schedule.

The initial schedule is generated with the use of a greedy heuristic. Firstly a resource utility rate u_i is calculated for each task

$$u_i = \frac{d_i \times r_{i,k}}{a_k} \tag{1}$$

For each task its resource requirement is divided by the resource availability. The result is multiplied by the task duration. Next all tasks are put into a list and ordered by non-increasing resource utility rate. After ordering each task is moved directly in front of the first successor in the list. The result is an activity list representation and the serial generation scheme is run on it to create the initial (left-justified) schedule.

After the initial list is generated the main iterative part of the algorithm starts with the destruction phase. During this phase a copy is made of the initial schedule and next $d = \lceil \frac{|N|}{4} \rceil$ tasks are removed from the activity list at random. These are picked one by one and are kept separately in the order they were removed.

The second step of the main iteration is the construction phase. From the removed tasks the first is picked and placed at any index in the remaining activity list that doesn't break the precedence order. For each possible index the makespan of the left-justified schedule made with the serial generation scheme is calculated. The index with the lowest makespan is

chosen and the task is inserted at the index. This process is repeated for each removed task until all tasks are in the activity list. At this point the resulting schedule makespan is compared to the initial schedule makespan and when an improved makespan has been found the initial schedule is overwritten by the new schedule.

This heuristic solution can any number of iterations of the destruction and construction phases until either a iterations limit is reached or a time limit is reached. At that point it will return the most optimal schedule it has found.

3.2 CNF Encoding

The CNF encoding is very much still in process and is therefore not a part of the draft version.

3.3 SAT solver

Due to difficulties during the CNF encoding my current planning does not allow for any time spent on improving the SAT solver.

4 Experimental Setup and Results

In order to test the performance of the different algorithmic approaches to solving PRCPSP-ST a number of experiments have been carried out. Each experiment is run on the high performance computing cluster at Delft University of Technology. The algorithms have access to 12GB RAM and 2 cores of the Intel(R) Xeon(R) CPU E5-26** v* running at 2.* GHz.

4.1 Project data

To test the difference between the approaches to include activity preemption a number of tests are performed using three different datasets. The datasets contain a different amount of problem instances (# instances) and project ranging from 10 to 50 tasks (# tasks). The J30 and RG30 datasets contain projects with 30 tasks and have 480 and 1800 instances respectively and the DC1 dataset has project ranging from 10 to 50 tasks also containing 1800 instances.

The setup time penalty s is set to 1, 2 and 5 to test the impact on the overall makespan. These values are chosen to be around .1, .2 and .5 of the longer tasks in the datasets that are around 10 time units in length.

As algorithms to test the tweaked version of the iterated greedy heuristic is used to calculate a baseline and the CNF encoding run on the pumpkin MAX-SAT solver is used to calculate data to compare to the baseline. Each algorithm is run for 10 (? not yet sure if this is in any way feasible and the limit might need to change an order of magnitude) seconds of CPU time on each instance.

To assess the performance two test values are calculated. First is the percentage deviation from the either the known optimal makespan or the lower bound makespan without activity preemption is calculated. This is used to check if both algorithms can come up with schedules that are reasonably close to existing solutions. The second test value is the percentage of makespans that is improved by the CNF endoing run on the pumpkin MAX-SAT solver compared to the iterated greedy heuristic solution. This measure will show if using a CNF encoding instead of a heuristic approach without

any further optimization could be used to find more optimal schedules.

4.2 Problem size

This subsection is to show the impact of the problem size on the performance of the algorithms. The experiments must still be run so this section remains empty for now.

4.3 Setup time

This subsection is to show the impact of the setup times on the resulting makespans when allowing preemption. The experiments must still be run so this section remains empty for now.

5 Responsible Research

In algorithmic optimization for NP-hard problems the impossibility of brute-force methods on large problem instances results in a competition for researchers to improve the current state-of-the-art solutions. Because a research is mostly focussed on positive results a number of criteria for optimizing the state of the art are often overlooked or ignored. The result is research papers that look like they indicate a more optimal algorithms has been found whereas the truth of the improvement might come from a different place altogether.

When technology evolves a lot of variables involved in algorithm testing are also improving, even without directly being a part of the algorithm. Examples are newer hardware, better compilers, improved coding skills and programming language differences. If a new algorithm is tested and these variables are not kept the same or the old versions are retested with the same new variables the resulting improvement might not be a result of a more optimal algorithm. Additionally, these variables do not cancel each other out. Over time it is to be expected that all the example variables are resulting in better performance when time passes.

To compare the performance of algorithms in a fair way some rules to keep variables equal are required. In this project the following rules have been set to keep the results more fair in comparison:

- Each algorithm is implemented in the same coding language
- The algorithms are implemented by the same researcher to ensure the coding skill is consistent
- Every algorithm is run on the same hardware
- Compiler and operating system are equal for all algorithms
- As a stopping criterion the CPU time is used

To make sure these rules can also be adhered to in the future the source code is provided to test the algorithms on newer hardware when it is released. Other researchers are also encouraged to implement the described algorithms for themselves and setup the experiments alongside their own algorithms.

6 Discussion

This section can only be written correctly after all experiments have been run and all results are evaluated.

7 Conclusions and Future Work

References

[1] W. Herroelen, E. Demeulemeester, and B. De Reyck, *A Classification Scheme for Project Scheduling*, pp. 1–26. Boston, MA: Springer US, 1999.

[2] P. Brucker, A. Drexl, R. Möhring, K. Neumann, and E. Pesch, "Resource-constrained project scheduling: Notation, classification, models, and methods," *European Journal of Operational Research*, vol. 112, no. 1, pp. 3–41, 1999.

[3] S. Hartmann and D. Briskorn, "A survey of variants and extensions of the resource-constrained project scheduling problem," *European Journal of Operational Research*, vol. 207, no. 1, pp. 1–14, 2010.

[4] E. Demeulemeester and W. Herroelen, "An efficient optimal solution procedure for the preemptive resource-constrained project scheduling problem," *European Journal of Operational Research*, vol. 90, no. 2, pp. 334–348, 1996.

[5] L. Kaplan, "Resource-constrained project scheduling with setup times." 1991.

[6] M. Vanhoucke and J. Coelho, "Resource-constrained project scheduling with activity splitting and setup times," *Computers & Operations Research*, vol. 109, pp. 230–249, 2019.

[7] J. Coelho and M. Vanhoucke, "Multi-mode resource-constrained project scheduling using rcpsp and sat solvers," *European Journal of Operational Research*, vol. 213, no. 1, pp. 73–82, 2011.

[8] D. Debels and M. Vanhoucke, "A decomposition-based genetic algorithm for the resource-constrained project-scheduling problem," *Operations Research*, vol. 55, no. 3, pp. 457–469, 2007.

[9] S. Hartmann, "A competitive genetic algorithm for resource-constrained project scheduling," *Naval Research Logistics (NRL)*, vol. 45, no. 7, pp. 733–750, 1998.

[10] J. Blazewicz, J. K. Lenstra, and A. H. G. Rinnooy Kan, "Scheduling subject to resource constraints: classification and complexity," *Discret. Appl. Math.*, vol. 5, pp. 11–24, 1983.

[11] S. Hartmann and D. Briskorn, "An updated survey of variants and extensions of the resource-constrained project scheduling problem," *European Journal of Operational Research*, vol. 297, no. 1, pp. 1–14, 2022.

[12] R. Ruiz and T. Stützle, "A simple and effective iterated greedy algorithm for the permutation flowshop scheduling problem," *European Journal of Operational Research*, vol. 177, no. 3, pp. 2033–2049, 2007.

[13] J. E. Kelley, "The critical-path method : Resource planning and scheduling," *Industrial Scheduling*, 1963.